

# Classification of MRI images using CNN for tumour detection

## (Part 1 - Running the models)


We will start by downloading the dataset from the Github repository and unzipping the contents of the repository into our local directory.

```
!wget https://github.com/SartajBhuvaji/Brain-Tumor-Classification-DataSet/archive/refs/heads/master.zip
!unzip -q master.zip
```

```
--2022-02-15 00:48:37-- https://github.com/SartajBhuvaji/Brain-Tumor-Classification-DataSet
Resolving github.com (github.com)... 140.82.114.4
Connecting to github.com (github.com)|140.82.114.4|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://codeload.github.com/SartajBhuvaji/Brain-Tumor-Classification-DataSet/archive/refs/heads/master.zip
--2022-02-15 00:48:37-- https://codeload.github.com/SartajBhuvaji/Brain-Tumor-Classification-DataSet/archive/refs/heads/master.zip
Resolving codeload.github.com (codeload.github.com)... 140.82.114.9
Connecting to codeload.github.com (codeload.github.com)|140.82.114.9|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: unspecified [application/zip]
Saving to: 'master.zip'
```

```
master.zip          [          <=> ]  86.97M  25.9MB/s   in 3.4s
```

```
2022-02-15 00:48:41 (25.3 MB/s) - 'master.zip' saved [91198591]
```



Let us now set the path to the directories containing the training and the validation images.

```
import pandas as pd
TRAIN_PATH = "./Brain-Tumor-Classification-DataSet-master/Training/"
TEST_PATH = "./Brain-Tumor-Classification-DataSet-master/Testing/"
```

We will now set the batch and the image size for the models.

```
import tensorflow as tf
batch_size = 32
img_height = 256
img_width = 256
```

We will now create the input pipeline Tensorflow-compatible dataset using the images in the training directory.

```
training_dataset = tf.keras.utils.image_dataset_from_directory(
    TRAIN_PATH,
    seed=0,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

Found 2870 files belonging to 4 classes.

Let us similarly create a dataset using the images in the validation directory.

```
testing_dataset = tf.keras.utils.image_dataset_from_directory(
    TEST_PATH,
    seed=0,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

Found 394 files belonging to 4 classes.

Let us view the names of the classes in the dataset(s).


```
classes = training_dataset.class_names
print(classes)
```

```
['glioma_tumor', 'meningioma_tumor', 'no_tumor', 'pituitary_tumor']
```

For our own visualization, let us plot a few of the images, along with their labels, from the dataset.

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in training_dataset.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(classes[labels[i]])
        plt.axis("off")
```



Let us configure Tensorflow to enable buffering for efficiency.

```

AUTOTUNE = tf.data.AUTOTUNE
training_dataset = training_dataset.cache().prefetch(buffer_size=AUTOTUNE)
testing_dataset = testing_dataset.cache().prefetch(buffer_size=AUTOTUNE)

```



Let us define a callback for early stopping so that the training can be automatically stopped in case there is no decrease in validation loss.

We will add a tolerance of five (5) epochs so that the training does not terminate in case the model is stuck in a plateau.

```

es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=

```

Let us create a variable to indicate the maximum number of epochs allowed for training the models.

```

MAX_EPOCHS = 100

```

## ▼ Model 1 - Basic Model

Let us define and build the basic model

```

basic_model = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255),
    tf.keras.layers.Conv2D(32, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(32, 3, activation='relu'),

```

```

tf.keras.layers.MaxPooling2D(),
tf.keras.layers.Conv2D(32, 3, activation='relu'),
tf.keras.layers.MaxPooling2D(),
tf.keras.layers.Flatten(),
tf.keras.layers.Dense(128, activation='relu'),
tf.keras.layers.Dense(4)
])

```

```

basic_model.build(input_shape=(None,256,256,3))
basic_model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 256, 256, 3)	0
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_2 (Conv2D)	(None, 60, 60, 32)	9248
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 32)	0
flatten (Flatten)	(None, 28800)	0
dense (Dense)	(None, 128)	3686528
dense_1 (Dense)	(None, 4)	516
Total params: 3,706,436		
Trainable params: 3,706,436		
Non-trainable params: 0		

We will now compile and train the model.

```

basic_model.compile(
    optimizer='adam',
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])

```

```

basic_history = basic_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,

```

```
callbacks=[es]
)
```

```
Epoch 1/100
90/90 [=====] - 21s 117ms/step - loss: 0.8842 - accuracy: 0.75
Epoch 2/100
90/90 [=====] - 8s 84ms/step - loss: 0.4886 - accuracy: 0.75
Epoch 3/100
90/90 [=====] - 8s 85ms/step - loss: 0.2971 - accuracy: 0.87
Epoch 4/100
90/90 [=====] - 8s 84ms/step - loss: 0.1683 - accuracy: 0.93
Epoch 5/100
90/90 [=====] - 8s 86ms/step - loss: 0.1074 - accuracy: 0.96
Epoch 6/100
90/90 [=====] - 8s 84ms/step - loss: 0.0642 - accuracy: 0.97
Epoch 7/100
90/90 [=====] - 8s 84ms/step - loss: 0.0407 - accuracy: 0.98
Epoch 8/100
90/90 [=====] - 8s 85ms/step - loss: 0.0235 - accuracy: 0.99
Epoch 9/100
90/90 [=====] - 8s 85ms/step - loss: 0.0089 - accuracy: 0.99
Epoch 10/100
90/90 [=====] - 8s 85ms/step - loss: 0.0073 - accuracy: 0.99
Epoch 11/100
90/90 [=====] - 8s 85ms/step - loss: 0.0096 - accuracy: 0.99
Epoch 12/100
90/90 [=====] - 8s 84ms/step - loss: 0.0078 - accuracy: 0.99
Epoch 00012: early stopping
```

Save the model and training history

```
basic_model.save("/content/drive/MyDrive/mri-cnn/basic")
pd.DataFrame(basic_history.history).to_csv("/content/drive/MyDrive/mri-cnn/basic.csv")
```

```
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/basic/assets
```

## ▼ Model 2 - ResNet50

Let us import the ResNet50 model.

```
resnet50_model = tf.keras.applications.ResNet50(include_top=False, weights="imagenet", inp
resnet50_model.summary()
```

```
conv5_block1_out (Activation) (None, 8, 8, 2048) 0 ['conv5_block1_ad
conv5_block2_1_conv (Conv2D) (None, 8, 8, 512) 1049088 ['conv5_block1_out
conv5_block2_1_bn (BatchNormal (None, 8, 8, 512) 2048 ['conv5_block2_1_
ization)
conv5_block2_1_relu (Activatio (None, 8, 8, 512) 0 ['conv5_block2_1_
n)
```

conv5_block2_2_conv (Conv2D)	(None, 8, 8, 512)	2359808	['conv5_block2_1_
conv5_block2_2_bn (BatchNormalization)	(None, 8, 8, 512)	2048	['conv5_block2_2_
conv5_block2_2_relu (Activation)	(None, 8, 8, 512)	0	['conv5_block2_2_
conv5_block2_3_conv (Conv2D)	(None, 8, 8, 2048)	1050624	['conv5_block2_2_
conv5_block2_3_bn (BatchNormalization)	(None, 8, 8, 2048)	8192	['conv5_block2_3_
conv5_block2_add (Add)	(None, 8, 8, 2048)	0	['conv5_block1_out', 'conv5_block2_3_
conv5_block2_out (Activation)	(None, 8, 8, 2048)	0	['conv5_block2_add
conv5_block3_1_conv (Conv2D)	(None, 8, 8, 512)	1049088	['conv5_block2_out
conv5_block3_1_bn (BatchNormalization)	(None, 8, 8, 512)	2048	['conv5_block3_1_
conv5_block3_1_relu (Activation)	(None, 8, 8, 512)	0	['conv5_block3_1_
conv5_block3_2_conv (Conv2D)	(None, 8, 8, 512)	2359808	['conv5_block3_1_
conv5_block3_2_bn (BatchNormalization)	(None, 8, 8, 512)	2048	['conv5_block3_2_
conv5_block3_2_relu (Activation)	(None, 8, 8, 512)	0	['conv5_block3_2_
conv5_block3_3_conv (Conv2D)	(None, 8, 8, 2048)	1050624	['conv5_block3_2_
conv5_block3_3_bn (BatchNormalization)	(None, 8, 8, 2048)	8192	['conv5_block3_3_
conv5_block3_add (Add)	(None, 8, 8, 2048)	0	['conv5_block2_out', 'conv5_block3_3_
conv5_block3_out (Activation)	(None, 8, 8, 2048)	0	['conv5_block3_add
=====			
Total params: 23,587,712			

We will flatten the final layer and then add our output layer with four (4) nodes.

```
res_model = tf.keras.models.Sequential()
res_model.add(resnet50_model)
res_model.add(tf.keras.layers.Flatten())
res_model.add(tf.keras.layers.Dense(4))
res_model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 8, 8, 2048)	23587712
flatten_1 (Flatten)	(None, 131072)	0
dense_2 (Dense)	(None, 4)	524292
Total params: 24,112,004		
Trainable params: 24,058,884		
Non-trainable params: 53,120		

We can finally compile and train the ResNet50 model.

```
res_model.compile(optimizer='adam',
                  loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])
res_history = res_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,
    callbacks=[es]
)
```

```
Epoch 1/100
90/90 [=====] - 88s 871ms/step - loss: 4.7750 - accuracy: 0
Epoch 2/100
90/90 [=====] - 74s 824ms/step - loss: 2.5800 - accuracy: 0
Epoch 3/100
90/90 [=====] - 74s 821ms/step - loss: 1.3020 - accuracy: 0
Epoch 4/100
90/90 [=====] - 74s 823ms/step - loss: 0.8957 - accuracy: 0
Epoch 5/100
90/90 [=====] - 75s 828ms/step - loss: 0.5563 - accuracy: 0
Epoch 6/100
90/90 [=====] - 75s 829ms/step - loss: 0.3237 - accuracy: 0
Epoch 7/100
90/90 [=====] - 75s 831ms/step - loss: 0.3140 - accuracy: 0
Epoch 8/100
90/90 [=====] - 75s 831ms/step - loss: 0.7986 - accuracy: 0
Epoch 9/100
90/90 [=====] - 75s 833ms/step - loss: 0.9672 - accuracy: 0
Epoch 10/100
90/90 [=====] - 75s 830ms/step - loss: 0.7650 - accuracy: 0
Epoch 11/100
90/90 [=====] - 75s 830ms/step - loss: 0.2536 - accuracy: 0
Epoch 12/100
90/90 [=====] - 75s 830ms/step - loss: 0.3117 - accuracy: 0
Epoch 13/100
90/90 [=====] - 75s 830ms/step - loss: 0.1172 - accuracy: 0
Epoch 00013: early stopping
```

## Save the model and training history

```
res_model.save("/content/drive/MyDrive/mri-cnn/res")
pd.DataFrame(res_history.history).to_csv("/content/drive/MyDrive/mri-cnn/res.csv")

INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/res/assets
/usr/local/lib/python3.7/dist-packages/keras/engine/functional.py:1410: CustomMaskWar
    layer_config = serialize_layer_fn(layer)
/usr/local/lib/python3.7/dist-packages/keras/saving/saved_model/layer_serialization.py:
    return generic_utils.serialize_keras_object(obj)
```

## ▼ Model 3 - VGG16

Let us import the VGG16 model.

```
vgg16_model = tf.keras.applications.VGG16(include_top=False, weights="imagenet", input_tensor=input_tensor)
vgg16_model.summary()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16\_weights\_tf\_dim\_ordering\_tf\_data\_format.h5
58892288/58889256 [=====] - 1s 0us/step
58900480/58889256 [=====] - 1s 0us/step
Model: "vgg16"
```

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0




block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0

```

=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0

```



Like before, let us add our output layers.

```

vgg_model = tf.keras.models.Sequential()
vgg_model.add(vgg16_model)
vgg_model.add(tf.keras.layers.Flatten())
vgg_model.add(tf.keras.layers.Dense(4))
vgg_model.summary()

```

```

Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 8, 8, 512)	14714688
flatten_2 (Flatten)	(None, 32768)	0
dense_3 (Dense)	(None, 4)	131076

```

=====
Total params: 14,845,764
Trainable params: 14,845,764
Non-trainable params: 0

```

We are now ready to compile and train the network.

```

vgg_model.compile(optimizer='adam',
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
vgg_history = vgg_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,
    callbacks=[es]
)

```

```

Epoch 1/100
90/90 [=====] - 127s 1s/step - loss: 2.7705 - accuracy: 0.36

```

```

Epoch 2/100
90/90 [=====] - 91s 1s/step - loss: 1.0757 - accuracy: 0.528
Epoch 3/100
90/90 [=====] - 91s 1s/step - loss: 0.9219 - accuracy: 0.616
Epoch 4/100
90/90 [=====] - 91s 1s/step - loss: 0.8214 - accuracy: 0.663
Epoch 5/100
90/90 [=====] - 90s 1s/step - loss: 0.6688 - accuracy: 0.743
Epoch 6/100
90/90 [=====] - 90s 1s/step - loss: 0.5753 - accuracy: 0.776
Epoch 7/100
90/90 [=====] - 90s 1s/step - loss: 0.5617 - accuracy: 0.776
Epoch 8/100
90/90 [=====] - 90s 1s/step - loss: 0.5124 - accuracy: 0.802
Epoch 9/100
90/90 [=====] - 90s 1s/step - loss: 0.4417 - accuracy: 0.832
Epoch 10/100
90/90 [=====] - 90s 1s/step - loss: 0.3813 - accuracy: 0.853
Epoch 11/100
90/90 [=====] - 90s 999ms/step - loss: 0.3649 - accuracy: 0
Epoch 12/100
90/90 [=====] - 90s 999ms/step - loss: 0.3138 - accuracy: 0
Epoch 00012: early stopping

```

Save the model and training history

```

vgg_model.save("/content/drive/MyDrive/mri-cnn/vgg")
pd.DataFrame(vgg_history.history).to_csv("/content/drive/MyDrive/mri-cnn/vgg.csv")

```

```
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/vgg/assets
```

## ▼ Model 4 - MobileNet

We will now import the MobileNet architecture.

```

mobilenet_model = tf.keras.applications.MobileNet(include_top=False, weights="imagenet", input_shape=(224, 224, 3))
mobilenet_model.summary()

```

conv_pw_10 (Conv2D)	(None, 16, 16, 512)	262144
conv_pw_10_bn (BatchNormalization)	(None, 16, 16, 512)	2048
conv_pw_10_relu (ReLU)	(None, 16, 16, 512)	0
conv_dw_11 (DepthwiseConv2D)	(None, 16, 16, 512)	4608
conv_dw_11_bn (BatchNormalization)	(None, 16, 16, 512)	2048
conv_dw_11_relu (ReLU)	(None, 16, 16, 512)	0
conv_pw_11 (Conv2D)	(None, 16, 16, 512)	262144

conv_pw_11_bn (BatchNormalization)	(None, 16, 16, 512)	2048
conv_pw_11_relu (ReLU)	(None, 16, 16, 512)	0
conv_pad_12 (ZeroPadding2D)	(None, 17, 17, 512)	0
conv_dw_12 (DepthwiseConv2D)	(None, 8, 8, 512)	4608
conv_dw_12_bn (BatchNormalization)	(None, 8, 8, 512)	2048
conv_dw_12_relu (ReLU)	(None, 8, 8, 512)	0
conv_pw_12 (Conv2D)	(None, 8, 8, 1024)	524288
conv_pw_12_bn (BatchNormalization)	(None, 8, 8, 1024)	4096
conv_pw_12_relu (ReLU)	(None, 8, 8, 1024)	0
conv_dw_13 (DepthwiseConv2D)	(None, 8, 8, 1024)	9216
conv_dw_13_bn (BatchNormalization)	(None, 8, 8, 1024)	4096
conv_dw_13_relu (ReLU)	(None, 8, 8, 1024)	0
conv_pw_13 (Conv2D)	(None, 8, 8, 1024)	1048576
conv_pw_13_bn (BatchNormalization)	(None, 8, 8, 1024)	4096
conv_pw_13_relu (ReLU)	(None, 8, 8, 1024)	0
=====		
Total params: 3,228,864		
Trainable params: 3,206,976		

Let us put our application-specific output layer into the imported architecture.

```
mobile_model = tf.keras.models.Sequential()
mobile_model.add(mobilenet_model)
mobile_model.add(tf.keras.layers.Flatten())
mobile_model.add(tf.keras.layers.Dense(4))
mobile_model.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
mobilenet_1.00_224 (Functional)	(None, 8, 8, 1024)	3228864

flatten_3 (Flatten)	(None, 65536)	0
dense_4 (Dense)	(None, 4)	262148

=====

Total params: 3,491,012  
Trainable params: 3,469,124  
Non-trainable params: 21,888

---

We can now compile and train our network.

```
mobile_model.compile(optimizer='adam',
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
mobile_history = mobile_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,
    callbacks=[es]
)
```

```
Epoch 1/100
90/90 [=====] - 42s 427ms/step - loss: 3.4591 - accuracy: 0
Epoch 2/100
90/90 [=====] - 37s 409ms/step - loss: 0.3927 - accuracy: 0
Epoch 3/100
90/90 [=====] - 37s 408ms/step - loss: 0.1201 - accuracy: 0
Epoch 4/100
90/90 [=====] - 37s 409ms/step - loss: 0.0469 - accuracy: 0
Epoch 5/100
90/90 [=====] - 37s 407ms/step - loss: 0.0358 - accuracy: 0
Epoch 6/100
90/90 [=====] - 37s 409ms/step - loss: 0.2890 - accuracy: 0
Epoch 7/100
90/90 [=====] - 37s 410ms/step - loss: 0.4403 - accuracy: 0
Epoch 8/100
90/90 [=====] - 37s 407ms/step - loss: 0.4581 - accuracy: 0
Epoch 9/100
90/90 [=====] - 37s 409ms/step - loss: 0.0878 - accuracy: 0
Epoch 10/100
90/90 [=====] - 37s 409ms/step - loss: 0.0662 - accuracy: 0
Epoch 11/100
90/90 [=====] - 37s 409ms/step - loss: 0.0534 - accuracy: 0
Epoch 12/100
90/90 [=====] - 37s 409ms/step - loss: 0.0249 - accuracy: 0
Epoch 13/100
90/90 [=====] - 37s 408ms/step - loss: 0.0068 - accuracy: 0
Epoch 14/100
90/90 [=====] - 37s 408ms/step - loss: 0.0104 - accuracy: 0
Epoch 15/100
90/90 [=====] - 37s 412ms/step - loss: 0.0499 - accuracy: 0
Epoch 16/100
90/90 [=====] - 37s 409ms/step - loss: 0.2978 - accuracy: 0
Epoch 17/100
90/90 [=====] - 37s 409ms/step - loss: 0.2518 - accuracy: 0
Epoch 18/100
```

```

90/90 [=====] - 37s 408ms/step - loss: 0.2289 - accuracy: 0
Epoch 19/100
90/90 [=====] - 37s 408ms/step - loss: 0.0580 - accuracy: 0
Epoch 20/100
90/90 [=====] - 37s 407ms/step - loss: 0.0428 - accuracy: 0
Epoch 21/100
90/90 [=====] - 37s 407ms/step - loss: 0.0124 - accuracy: 0
Epoch 22/100
90/90 [=====] - 37s 408ms/step - loss: 0.0136 - accuracy: 0
Epoch 00022: early stopping

```

## Save the model and training history

```

mobile_model.save("/content/drive/MyDrive/mri-cnn/mobile")
pd.DataFrame(mobile_history.history).to_csv("/content/drive/MyDrive/mri-cnn/mobile.csv")

WARNING:absl:Function `_wrapped_model` contains input name(s) mobilenet_1.00_224_inp
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/mobile/assets
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/mobile/assets

```

## ▼ Model 5 - Inception v3

Let us load the InceptionNetv3 model.

```

inceptionnet_model = tf.keras.applications.InceptionV3(include_top=False, weights="imagenet")
inceptionnet_model.summary()

```

conv2d_90 (Conv2D)	(None, 6, 6, 384)	442368	['activation_86[0]
conv2d_91 (Conv2D)	(None, 6, 6, 384)	442368	['activation_86[0]
conv2d_94 (Conv2D)	(None, 6, 6, 384)	442368	['activation_90[0]
conv2d_95 (Conv2D)	(None, 6, 6, 384)	442368	['activation_90[0]
average_pooling2d_8 (AveragePooling2D)	(None, 6, 6, 2048)	0	['mixed9[0][0]']
conv2d_88 (Conv2D)	(None, 6, 6, 320)	655360	['mixed9[0][0]']
batch_normalization_87 (Batch Normalization)	(None, 6, 6, 384)	1152	['conv2d_90[0][0]
batch_normalization_88 (Batch Normalization)	(None, 6, 6, 384)	1152	['conv2d_91[0][0]
batch_normalization_91 (Batch Normalization)	(None, 6, 6, 384)	1152	['conv2d_94[0][0]
batch_normalization_92 (Batch Normalization)	(None, 6, 6, 384)	1152	['conv2d_95[0][0]
conv2d_96 (Conv2D)	(None, 6, 6, 192)	393216	['average_pooling

conv2d_90 (Conv2D)	(None, 6, 6, 192)	55216	['average_pooling_84[0][0]']
batch_normalization_85 (Batch Normalization)	(None, 6, 6, 320)	960	['conv2d_88[0][0]']
activation_87 (Activation)	(None, 6, 6, 384)	0	['batch_normalization_85[0][0]']
activation_88 (Activation)	(None, 6, 6, 384)	0	['batch_normalization_85[0][0]']
activation_91 (Activation)	(None, 6, 6, 384)	0	['batch_normalization_85[0][0]']
activation_92 (Activation)	(None, 6, 6, 384)	0	['batch_normalization_85[0][0]']
batch_normalization_93 (Batch Normalization)	(None, 6, 6, 192)	576	['conv2d_96[0][0]']
activation_85 (Activation)	(None, 6, 6, 320)	0	['batch_normalization_93[0][0]']
mixed9_1 (Concatenate)	(None, 6, 6, 768)	0	['activation_87[0][0]', 'activation_88[0][0]']
concatenate_1 (Concatenate)	(None, 6, 6, 768)	0	['activation_91[0][0]', 'activation_92[0][0]']
activation_93 (Activation)	(None, 6, 6, 192)	0	['batch_normalization_93[0][0]']
mixed10 (Concatenate)	(None, 6, 6, 2048)	0	['activation_85[0][0]', 'mixed9_1[0][0]', 'concatenate_1[0][0]', 'activation_93[0][0]']

As before, we will add our application-specific output layer with four (4) nodes.

```
inception_model = tf.keras.models.Sequential()
inception_model.add(inceptionnet_model)
inception_model.add(tf.keras.layers.Flatten())
inception_model.add(tf.keras.layers.Dense(4))
inception_model.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
inception_v3 (Functional)	(None, 6, 6, 2048)	21802784
flatten_4 (Flatten)	(None, 73728)	0
dense_5 (Dense)	(None, 4)	294916

=====  
 Total params: 22,097,700  
 Trainable params: 22,063,268  
 Non-trainable params: 34,432

We are ready to compile and train the neural network.

```
inception_model.compile(optimizer='adam',
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
inception_history = inception_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,
    callbacks=[es]
)
```

```
Epoch 1/100
90/90 [=====] - 86s 821ms/step - loss: 2.1325 - accuracy: 0
Epoch 2/100
90/90 [=====] - 67s 746ms/step - loss: 2.7610 - accuracy: 0
Epoch 3/100
90/90 [=====] - 67s 745ms/step - loss: 0.8254 - accuracy: 0
Epoch 4/100
90/90 [=====] - 67s 747ms/step - loss: 0.6639 - accuracy: 0
Epoch 5/100
90/90 [=====] - 67s 744ms/step - loss: 0.4724 - accuracy: 0
Epoch 6/100
90/90 [=====] - 67s 745ms/step - loss: 0.3130 - accuracy: 0
Epoch 7/100
90/90 [=====] - 67s 747ms/step - loss: 0.1806 - accuracy: 0
Epoch 8/100
90/90 [=====] - 67s 747ms/step - loss: 0.1599 - accuracy: 0
Epoch 9/100
90/90 [=====] - 67s 745ms/step - loss: 0.1022 - accuracy: 0
Epoch 10/100
90/90 [=====] - 67s 748ms/step - loss: 0.0704 - accuracy: 0
Epoch 11/100
90/90 [=====] - 67s 746ms/step - loss: 0.0675 - accuracy: 0
Epoch 12/100
90/90 [=====] - 67s 746ms/step - loss: 0.0752 - accuracy: 0
Epoch 13/100
90/90 [=====] - 67s 746ms/step - loss: 0.0675 - accuracy: 0
Epoch 14/100
90/90 [=====] - 67s 748ms/step - loss: 0.0484 - accuracy: 0
Epoch 15/100
90/90 [=====] - 67s 748ms/step - loss: 0.0552 - accuracy: 0
Epoch 16/100
90/90 [=====] - 67s 748ms/step - loss: 0.0590 - accuracy: 0
Epoch 17/100
90/90 [=====] - 67s 746ms/step - loss: 0.0286 - accuracy: 0
Epoch 00017: early stopping
```



Save the model and training history

```
inception_model.save("/content/drive/MyDrive/mri-cnn/inception")
pd.DataFrame(inception_history.history).to_csv("/content/drive/MyDrive/mri-cnn/inception.c
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/inception/assets
```

## ▼ Model 6 - AlexNet

The Tensorflow core applications have no prebuilt model for AlexNet.

We will build the model on our own.

```
alexnet_model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    tf.keras.layers.Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    tf.keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(4096, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(4096, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

```
alexnet_model.build(input_shape=(None,256,256,3))
alexnet_model.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_97 (Conv2D)	(None, 62, 62, 96)	34944
batch_normalization_94 (Batch Normalization)	(None, 62, 62, 96)	384
max_pooling2d_7 (MaxPooling2D)	(None, 30, 30, 96)	0
conv2d_98 (Conv2D)	(None, 30, 30, 256)	614656
batch_normalization_95 (Batch Normalization)	(None, 30, 30, 256)	1024
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 256)	0
conv2d_99 (Conv2D)	(None, 14, 14, 384)	885120



batch_normalization_96 (Batch Normalization)	(None, 14, 14, 384)	1536
conv2d_100 (Conv2D)	(None, 14, 14, 384)	1327488
batch_normalization_97 (Batch Normalization)	(None, 14, 14, 384)	1536
conv2d_101 (Conv2D)	(None, 14, 14, 256)	884992
batch_normalization_98 (Batch Normalization)	(None, 14, 14, 256)	1024
max_pooling2d_9 (MaxPooling2D)	(None, 6, 6, 256)	0
flatten_5 (Flatten)	(None, 9216)	0
dense_6 (Dense)	(None, 4096)	37752832
dropout (Dropout)	(None, 4096)	0
dense_7 (Dense)	(None, 4096)	16781312
dropout_1 (Dropout)	(None, 4096)	0
dense_8 (Dense)	(None, 10)	40970
=====		
Total params: 58,327,818		
Trainable params: 58,325,066		
Non-trainable params: 2,752		

---

Let us compile and train the AlexNet model.

```
alexnet_model.compile(optimizer='adam',
                      loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
                      metrics=['accuracy'])
alexnet_history = alexnet_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,
    callbacks=[es]
)
```

```
Epoch 5/100
90/90 [=====] - 11s 126ms/step - loss: 1.1048 - accuracy:
Epoch 6/100
90/90 [=====] - 11s 126ms/step - loss: 1.1182 - accuracy:
Epoch 7/100
90/90 [=====] - 11s 126ms/step - loss: 1.0511 - accuracy:
Epoch 8/100
90/90 [=====] - 12s 128ms/step - loss: 0.9755 - accuracy:
Epoch 9/100
90/90 [=====] - 11s 127ms/step - loss: 0.9343 - accuracy:
Epoch 10/100
90/90 [=====] - 11s 127ms/step - loss: 0.9337 - accuracy:
```

```

90/90 [=====] - 11s 127ms/step - loss: 0.9387 - accuracy:
Epoch 11/100
90/90 [=====] - 12s 128ms/step - loss: 0.9197 - accuracy:
Epoch 12/100
90/90 [=====] - 11s 127ms/step - loss: 0.8904 - accuracy:
Epoch 13/100
90/90 [=====] - 11s 127ms/step - loss: 0.7953 - accuracy:
Epoch 14/100
90/90 [=====] - 11s 127ms/step - loss: 0.8727 - accuracy:
Epoch 15/100
90/90 [=====] - 11s 127ms/step - loss: 0.7750 - accuracy:
Epoch 16/100
90/90 [=====] - 11s 127ms/step - loss: 0.7289 - accuracy:
Epoch 17/100
90/90 [=====] - 11s 127ms/step - loss: 0.7285 - accuracy:
Epoch 18/100
90/90 [=====] - 11s 127ms/step - loss: 0.7193 - accuracy:
Epoch 19/100
90/90 [=====] - 11s 127ms/step - loss: 0.7022 - accuracy:
Epoch 20/100
90/90 [=====] - 11s 127ms/step - loss: 0.7323 - accuracy:
Epoch 21/100
90/90 [=====] - 11s 127ms/step - loss: 0.6812 - accuracy:
Epoch 22/100
90/90 [=====] - 11s 127ms/step - loss: 0.5751 - accuracy:
Epoch 23/100
90/90 [=====] - 11s 127ms/step - loss: 0.5924 - accuracy:
Epoch 24/100
90/90 [=====] - 11s 127ms/step - loss: 0.5515 - accuracy:
Epoch 25/100
90/90 [=====] - 11s 127ms/step - loss: 0.6511 - accuracy:
Epoch 26/100
90/90 [=====] - 12s 128ms/step - loss: 0.5781 - accuracy:
Epoch 27/100
90/90 [=====] - 11s 127ms/step - loss: 0.5879 - accuracy:
Epoch 28/100
90/90 [=====] - 11s 126ms/step - loss: 0.5588 - accuracy:
Epoch 29/100
90/90 [=====] - 11s 127ms/step - loss: 0.5459 - accuracy:
Epoch 30/100
90/90 [=====] - 12s 127ms/step - loss: 0.4622 - accuracy:
Epoch 31/100
90/90 [=====] - 11s 126ms/step - loss: 0.5897 - accuracy:
Epoch 32/100
90/90 [=====] - 11s 127ms/step - loss: 0.4947 - accuracy:
Epoch 00032: early stopping

```

## Save the model and training history

```

alexnet_model.save("/content/drive/MyDrive/mri-cnn/alexnet")
pd.DataFrame(alexnet_history.history).to_csv("/content/drive/MyDrive/mri-cnn/alexnet.csv")

```

```

INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/alexnet/assets
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/alexnet/assets

```

## ▼ Model 7 - LeNet

Like AlexNet, Tensorflow does not have a prebuilt model for LeNet. Therefore, we will define the model on our own.

```
lenet_model = tf.keras.Sequential()
lenet_model.add(tf.keras.layers.Conv2D(filters=6, kernel_size=(3, 3), activation='relu', i
lenet_model.add(tf.keras.layers.AveragePooling2D())
lenet_model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=(3, 3), activation='relu'))
lenet_model.add(tf.keras.layers.AveragePooling2D())
lenet_model.add(tf.keras.layers.Flatten())
lenet_model.add(tf.keras.layers.Dense(units=120, activation='relu'))
lenet_model.add(tf.keras.layers.Dense(units=84, activation='relu'))
lenet_model.add(tf.keras.layers.Dense(units=4))
```

We are now ready to compile and train LeNet

```
lenet_model.compile(optimizer='adam',
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
lenet_history = lenet_model.fit(
    training_dataset,
    validation_data=testing_dataset,
    epochs=MAX_EPOCHS,
    callbacks=[es]
)
```

```
Epoch 1/100
90/90 [=====] - 5s 45ms/step - loss: 9.0646 - accuracy: 0.61
Epoch 2/100
90/90 [=====] - 4s 41ms/step - loss: 0.2646 - accuracy: 0.90
Epoch 3/100
90/90 [=====] - 4s 42ms/step - loss: 0.1433 - accuracy: 0.95
Epoch 4/100
90/90 [=====] - 4s 42ms/step - loss: 0.0894 - accuracy: 0.96
Epoch 5/100
90/90 [=====] - 4s 42ms/step - loss: 0.0548 - accuracy: 0.98
Epoch 6/100
90/90 [=====] - 4s 40ms/step - loss: 0.0087 - accuracy: 0.99
Epoch 7/100
90/90 [=====] - 4s 40ms/step - loss: 0.0040 - accuracy: 1.00
Epoch 8/100
90/90 [=====] - 4s 40ms/step - loss: 0.0011 - accuracy: 1.00
Epoch 9/100
90/90 [=====] - 4s 40ms/step - loss: 6.6158e-04 - accuracy:
Epoch 10/100
90/90 [=====] - 4s 40ms/step - loss: 5.1745e-04 - accuracy:
Epoch 11/100
90/90 [=====] - 4s 39ms/step - loss: 4.2063e-04 - accuracy:
Epoch 12/100
90/90 [=====] - 4s 40ms/step - loss: 3.4883e-04 - accuracy:
Epoch 13/100
90/90 [=====] - 4s 40ms/step - loss: 2.9345e-04 - accuracy:
Epoch 00013: early stopping
```

## Save the model and training history

```
lenet_model.save("/content/drive/MyDrive/mri-cnn/lenet")
pd.DataFrame(lenet_history.history).to_csv("/content/drive/MyDrive/mri-cnn/lenet.csv")

INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/lenet/assets
INFO:tensorflow:Assets written to: /content/drive/MyDrive/mri-cnn/lenet/assets
```

## Conclusion

We have used the following models for the MRI tumour classification task -

1. Basic CNN
2. ResNet50
3. VGG16
4. MobileNet
5. InceptionNet v2
6. AlexNet
7. LeNet

We splitted the dataset into training and validation sets. Once splitted, we trained the different models using the training set.

We used the validation set to determine the training performance at the end of each epoch.

We used early stopping with a high tolerance to stop the training earlier than the scheduled time in case the training was stuck in a minima without further improvement in performance.

We have saved the models trained and the training histories for further analysis.

